

# GAE Feedback

Distributed Systems 2015-2016

# Key Criteria

- 1. Persistence: GAE Datastore using JPA
  - Structure of entities in Entity Groups
  - Use of JPQL
- 2. Indirect Communication
  - Adapt from Direct Communication
  - Use of Task Queues, reason about data flow
- 3. Potential State Inconsistencies (theoretical exercise)
  - Why? How to avoid?
  - How to minimize impact on performance?

# General Impression

## ■ Observation:

- ✓ Persistence at GAE
- ! Indirect communication
- ! State inconsistency appearing in worker-based schema

## ■ Hotspots for erros


- JPA: attached vs. detached objects
- Indirection == Time Decoupling
- TaskQueue: Worker threads vs. processes

## ■ Written answers

- Some are quite good (short + complete), even providing alternative reasons
- Others focus too much on the application and less on concepts and technology (usually the (too) long ones)

# 1. Persistence

- **Entity Group (EG):** entity relations as hierarchical **tree-structure**
  - one-to-many == parent-child relationship
  - every entity has at max one parent entity (tree)
- **Easiest solution in car rental agency**
  - All entities in a single Entity Group
  - Use EntityManager instance per modification of EG
    - E.g. Multiple CRCs == multiple EntityManagers
  - No unmanaged relations between those entities
    - E.g. Manual lookup by UID (Primary Key)



# 1. Persistence (cont.)

- Transactions
  - Each EntityManager session == atomic commit
  - Don't use CrossGroup transactions (XGT)
- Google App Engine: JPQL
  - No JOINS supported
  - Use only to query on one entity

# 1. Persistence (cont.)

- JPA: Attached vs. Detached entities

- Change is not persisted:

```
EntityManager em = EMF.get().createEntityManager();  
List<Quote> q = em.createQuery(..).getResultList();  
em.close();  
[...]  
q.get(0).setStatus(Status.FAILED);
```


- Change is persisted:

```
EntityManager em = ...  
List<Quote> q = ...  
q.get(0).setStatus(Status.FAILED);  
em.close();
```

```
em.close();  
q.get(0).setStatus(Status.FAILED);  
em=EMF.get()...  
em.merge(q);
```

# 2. Indirect Communication

- Principle of Message Queues understood
- Confusion with the **3 roles** involved
  - Data flow:
    - Client (Browser)  $\leftarrow$  direct  $\rightarrow$  Font-end Service
    - Font-end Service – via Queue  $\rightarrow$  Back-end Service (Worker)



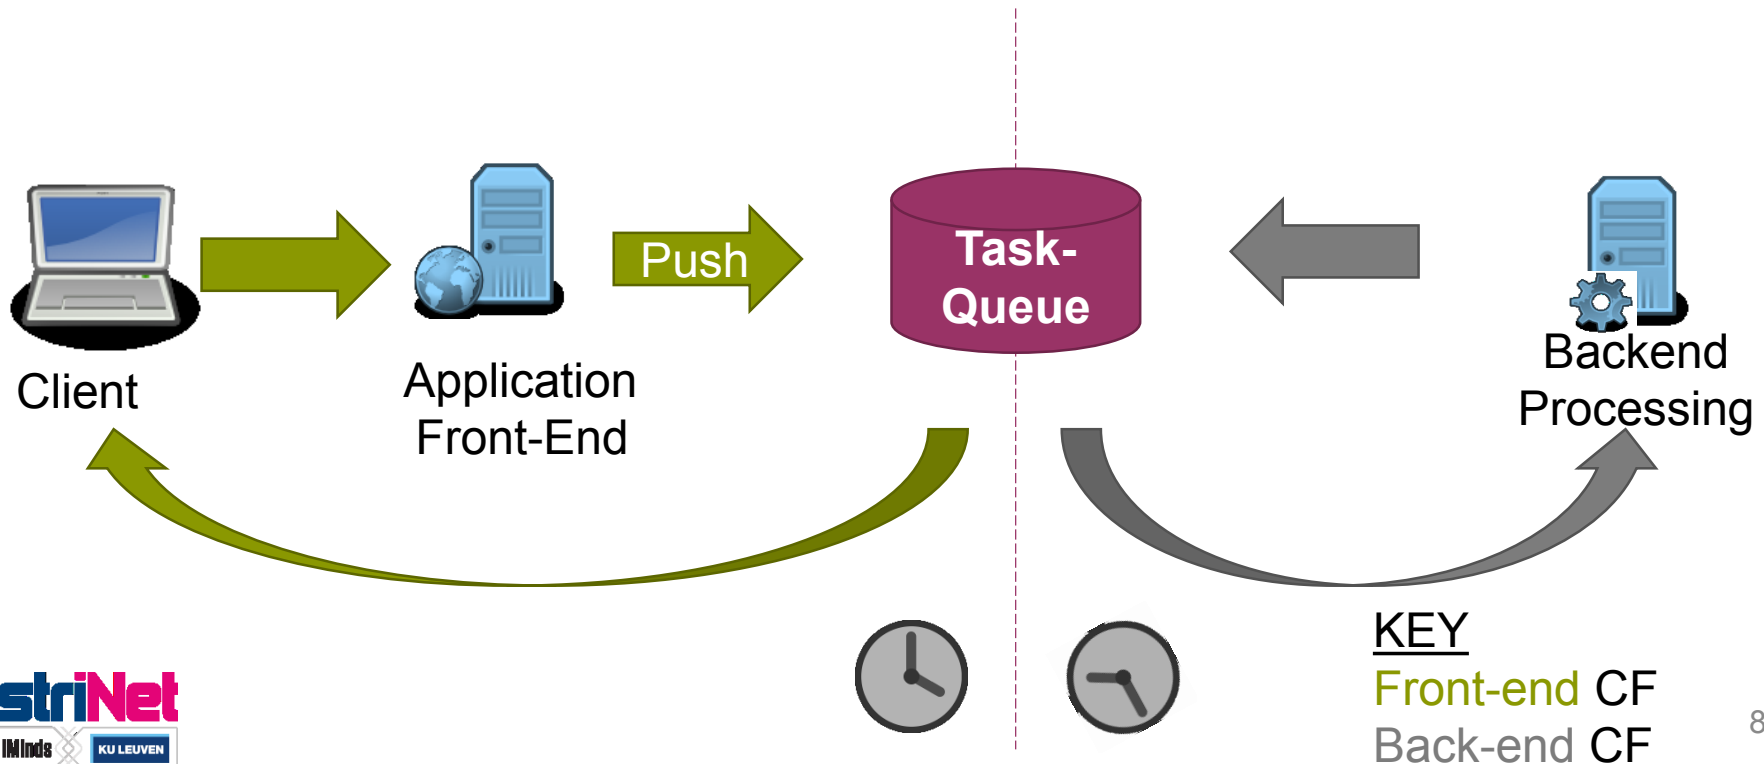
## KEY

Direct Communication  
Indirect Communication

# 2. Indirect Communication

- ...
- Control flow (CF):

- Front-end Service triggered by Client
- Back-end Service triggered by GAE infrastructure
- (1) and (2) are **independent in time**





# 2. Indirect Communication

- Do not
  - Send msgs from back-end to front-end
    - client may be gone at time of message
  - wait with front-end until back-end completes
- Examples
  - no shared memory (shared „global“ variables )
  - no Channel API
- Backend workers use separate machines
  - in own process ( $\Leftrightarrow$  own thread)
  - no shared memory
    - shared „global“ static variables > won't work

## 2. Indirect Communication

- **Feedback channel** from Worker to Client
  - report about success AND failure
  - otherwise indistinguishable : not-yet-processed or failed
- Client may have **multiple** unconfirmed **quotes waiting**, or may have sent a rental order twice.
  - feedback should uniquely refer to quotes (e.g. using **orderId**)

# 3. Potential State Inconsistencies

- **State changes** (creation of reservations) happen only at the Worker role and **not at the front-end**.
  - Keyword **‘synchronized’** for quote-enqueue function **no effect** on potential state inconsistencies.
- Multiple workers process tasks queue **in parallel** (default)
  - This is where potential inconsistencies are rooted
- Each worker is a separate process (separate JVM)
  - thread-based monitors (i.e. **‘synchronized’**) **cannot help**

# 3. Potential State Inconsistencies

## ■ One solution

- Set #worker per queue = 1
- To increase parallelism
  - Maintain one queue per company  
(**Note:** only if client will book from single CRC)

## ■ NOT

- Create a queue per CarType or per Quote
  - → loss of all-or-nothing semantics
- Create a queue for every task
  - jeopardizes state consistency even more (parallelism  $\rightarrow \infty$ )

# Written Answers

- Always: reason about your answer (unless it's certainly obvious)
  - Why did you choose tech a or option b?
  - Can you imagine drawbacks?
- Be to the point and don't loose the target
  - Asked: Is there a scenario for an inconsistent state.. ?
  - Expected: Yes, (the scenario is) when a does b.
  - Found: The application receives two requests [...] Suppose a sequence like [...] now suppose another sequence [...] As a result, [...] return OK, but [...] is the updated value.
  - Result:
    - after 1 page text
    - not sure whether you mean yes or no?

# Written Answers

- Another DO-NOT example: deferred answer
  - If there are multiple threads running to confirm quotes. There is a possibility that two consecutive tasks that are in conflict, when multiple threads are running they can confirm the quotes at the same time.
  - Read: if you believe that multiple threads are in use, than yes, otherwise no

# Questions?