

## Session 4

- Naive Bayes
- Bayesian Networks
- Reinforcement Learning

# Naive Bayes

- $V$  = set of classes
- $a_1, \dots, a_n$  = set of attributes
- $v = \operatorname{argmax}_{v_j \in V} P(v_j | a_1, a_2, \dots, a_n)$

- Apply Bayes rule

$$v = \operatorname{argmax}_{v_j \in V} \frac{P(a_1, a_2, \dots, a_n | v_j) \cdot P(v_j)}{P(a_1, a_2, \dots, a_n)}$$

- Remove constant factor

$$v = \operatorname{argmax}_{v_j \in V} P(a_1, a_2, \dots, a_n | v_j) \cdot P(v_j)$$

- Assume:

attribute values conditionally independent given class

$$v = \operatorname{argmax}_{v_j \in V} P(v_j) \cdot \prod_i P(a_i | v_j)$$

# Naive Bayes for Text Classification

**procedure** naive\_bayes\_text( $E$ : texts,  $V$ : classes, Voc: vocabulary)

for all texts  $A_i$  in  $E$ :

    preprocess( $A_i$ , Voc)

estimate  $P(v_j)$  and  $P(w_k|v_j)$  for all  $w_k$  in Voc and  $v_j$  in  $V$ :

$N_j$  = number of texts of class  $j$

$N$  = number of texts

$P(v_j) = N_j/N$

$n_{kj}$  = number of occurrences of  $w_k$  in texts of class  $j$

$n_j$  = number of words in class  $j$  (counting doubles)

$P(w_k|v_j) = (n_{kj} + 1)/(n_j + |\text{Voc}|)$

**procedure** classify\_naive\_bayes\_text( $A$ : text, Voc: vocabulary)

    preprocess( $A$ , Voc)

**return**  $\text{argmax}_{v_j \in V} P(v_j) \times \prod_{w_i \in \text{words}(A)} P(w_i|v_j)$

**procedure** preprocess( $A$ : text, Voc: vocabulary)

    transform each word in  $A$  to a standard form (stemming)

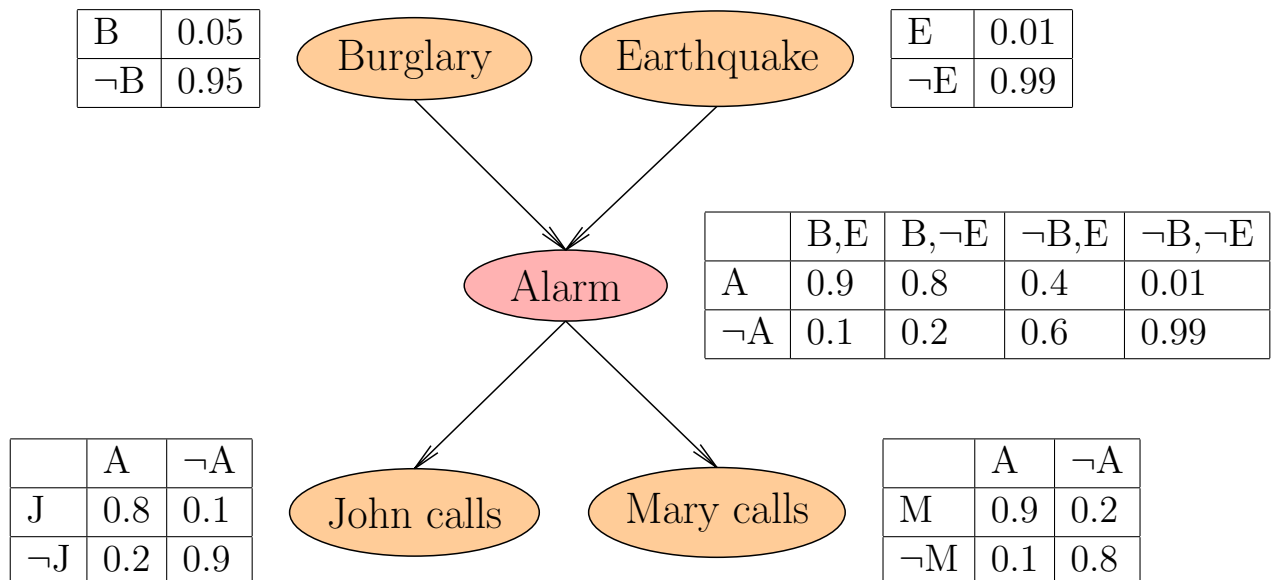
        e.g. *wait*, *waits*, *waited*, *waiting*  $\Rightarrow$  *wait*

    remove from  $A$  all words/tokens that are not in Voc


# Bayesian Networks

- Graph that represents joint probability distribution
- Node = attribute
- Edge = “direct influence” between two attributes
- Formally:

each node is conditionally independent of each of its non-descendants, given its parents



# Reinforcement Learning



- State space  $S$
- Set of possible actions  $A$
- Learn policy  $\pi : S \rightarrow A$
- Unknown reward function  $r : S \times A \rightarrow \mathbb{R}$
- Unknown next state function  $\delta : S \times A \rightarrow S$

Value function:  $V^\pi(s) = \sum_{t=0}^{\infty} \gamma^t \cdot r(s_t, a_t)$

Discounted future rewards from following policy  $\pi$  from state  $s$

Optimal policy:  $\pi^*(s) = \operatorname{argmax}_a \overbrace{r(s, a) + \gamma \cdot V^*(\delta(s, a))}^{Q(s, a)}$

Selects action  $a$  that maximises immediate reward and discounted future rewards

$$V^*(s) = V^{\pi^*}(s) = \max_a Q(s, a)$$

$$Q(s, a) = r(s, a) + \gamma \cdot \max_{a'} Q(\delta(s, a), a')$$

